

Algorithms Used for Classification

1. CART (Classification and Regression Trees)
2. Gaussian Naive Bayes / Naive Bayes
3. Gradient Boosting Machines (AdaBoost)
4. K-Nearest Neighbors (K-NN)
5. Logistic Regression
6. Multi-Layer Perceptron (MLP)
7. Perceptron
8. Random Forest
9. Support Vector Machines (SVM)

1. CART (Classification and Regression Trees) - DecisionTree Classifier

- Sampling Technique - Train/Test Split (80:20)
- Classification Metrics - Accuracy

```
from pandas import read_csv
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/car_evaluation_classification.csv'
dataframe = read_csv(filename)
# Custom mapping for each feature
buying_price_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
maintenance_cost_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
lug_boot_mapping = {'big': 2, 'med': 1, 'small': 0}
safety_mapping = {'high': 2, 'med': 1, 'low': 0}

# Apply the custom mapping to each column
dataframe['Buying Price'] = dataframe['Buying
Price'].map(buying_price_mapping)
dataframe['Maintenance Cost'] = dataframe['Maintenance
Cost'].map(maintenance_cost_mapping)
dataframe['Lug_Boot'] = dataframe['Lug_Boot'].map(lug_boot_mapping)
dataframe['Safety'] = dataframe['Safety'].map(safety_mapping)

# features (X) and target (Y)
X = dataframe[['Buying Price', 'Maintenance Cost', 'Number of Doors',
'Number of Persons', 'Lug_Boot', 'Safety']]
Y = dataframe['Classification']

# Set the test size and random seed for reproducibility
test_size = 0.20
random_seed = 50 # You can change this value

# Split the dataset into training and testing sets
```

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=test_size, random_state=random_seed)
```

```
# Initialize and train the Decision Tree Classifier with hyperparameters
```

```
max_depth = 5 # You can adjust this value
min_samples_split = 2 # You can adjust this value
min_samples_leaf = 1 # You can adjust this value
```

```
model = DecisionTreeClassifier(
    max_depth=max_depth,
    min_samples_split=min_samples_split,
    min_samples_leaf=min_samples_leaf,
    random_state=random_seed
)
```

```
model.fit(X_train, Y_train)
```

```
# Evaluate the accuracy
```

```
accuracy = model.score(X_test, Y_test)
print("Accuracy: %.3f%%" % (accuracy * 100.0))
```

```
Accuracy: 93.353%
```

2. Gaussian Naive Bayes

- Sampling Technique - Train/Test Split (80:20)
- Classification Metrics - Accuracy

```
from pandas import read_csv
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
```

```
# Load the dataset
```

```
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/car_evaluation_classification.csv'
dataframe = read_csv(filename)
```

```
# Custom mapping for each feature
```

```
buying_price_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
maintenance_cost_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
lug_boot_mapping = {'big': 2, 'med': 1, 'small': 0}
safety_mapping = {'high': 2, 'med': 1, 'low': 0}
```

```
# Apply the custom mapping to each column
```

```
dataframe['Buying Price'] = dataframe['Buying
Price'].map(buying_price_mapping)
dataframe['Maintenance Cost'] = dataframe['Maintenance
Cost'].map(maintenance_cost_mapping)
dataframe['Lug Boot'] = dataframe['Lug Boot'].map(lug_boot_mapping)
dataframe['Safety'] = dataframe['Safety'].map(safety_mapping)
```

```

# features (X) and target (Y)
X = dataframe[['Buying Price', 'Maintenance Cost', 'Number of Doors',
'Number of Persons', 'Lug_Boot', 'Safety']]
Y = dataframe['Classification']

# Set the test size
test_size = 0.20 # Hyperparameter: Fraction of the dataset to use for
testing
seed = 7

# Split the dataset into test and train
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=test_size, random_state=seed)

# Create a Gaussian Naive Bayes classifier
model = GaussianNB(priors=None, var_smoothing=1e-9)
# Hyperparameters:
# - priors: You can specify class prior probabilities if you have
prior knowledge.
# - var_smoothing: A smoothing parameter for avoiding zero variances.

# Train the model on the training data
model.fit(X_train, Y_train)

# Evaluate the accuracy
result = model.score(X_test, Y_test)
print("Accuracy: %.3f%%" % (result * 100.0))

Accuracy: 92.775%

```

3. Gradient Boosting Machines (AdaBoost)

- Sampling Technique - Train/Test Split (80:20)
- Classification Metrics - Accuracy

```

from pandas import read_csv
from sklearn.model_selection import train_test_split
from sklearn.ensemble import AdaBoostClassifier

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/car_evaluation_classification.csv'
dataframe = read_csv(filename)
# Custom mapping for each feature
buying_price_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
maintenance_cost_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
lug_boot_mapping = {'big': 2, 'med': 1, 'small': 0}
safety_mapping = {'high': 2, 'med': 1, 'low': 0}

# Apply the custom mapping to each column
dataframe['Buying Price'] = dataframe['Buying

```

```

Price'].map(buying_price_mapping)
dataframe['Maintenance Cost'] = dataframe['Maintenance
Cost'].map(maintenance_cost_mapping)
dataframe['Lug_Boot'] = dataframe['Lug_Boot'].map(lug_boot_mapping)
dataframe['Safety'] = dataframe['Safety'].map(safety_mapping)

# features (X) and target (Y)
X = dataframe[['Buying Price', 'Maintenance Cost', 'Number of Doors',
'Number of Persons', 'Lug_Boot', 'Safety']]
Y = dataframe['Classification']

# Set the test size
test_size = 0.20 # Hyperparameter: Fraction of the dataset to use for
testing
seed = 7

# Split the dataset into test and train
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=test_size, random_state=seed)

# Create an AdaBoost classifier
model = AdaBoostClassifier(n_estimators=50, random_state=seed)
# Hyperparameters:
# - n_estimators: The number of weak classifiers (base estimators) to
train. You can adjust this to control the complexity of the ensemble.
# - random_state: The random seed for reproducibility. You can set
this to a specific value if you want consistent results.

# Train the model on the training data
model.fit(X_train, Y_train)

# Evaluate the accuracy
result = model.score(X_test, Y_test)
print("Accuracy: %.3f%%" % (result * 100.0))

Accuracy: 95.087%

```

4. K-Nearest Neighbors (K-NN)

- Sampling Technique - Train/Test Split (80:20)
- Classification Metrics - Accuracy

```

from pandas import read_csv
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/car_evaluation_classification.csv'
dataframe = read_csv(filename)
# Custom mapping for each feature

```

```

buying_price_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
maintenance_cost_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
lug_boot_mapping = {'big': 2, 'med': 1, 'small': 0}
safety_mapping = {'high': 2, 'med': 1, 'low': 0}

# Apply the custom mapping to each column
dataframe['Buying Price'] = dataframe['Buying
Price'].map(buying_price_mapping)
dataframe['Maintenance Cost'] = dataframe['Maintenance
Cost'].map(maintenance_cost_mapping)
dataframe['Lug_Boot'] = dataframe['Lug_Boot'].map(lug_boot_mapping)
dataframe['Safety'] = dataframe['Safety'].map(safety_mapping)

# features (X) and target (Y)
X = dataframe[['Buying Price', 'Maintenance Cost', 'Number of Doors',
'Number of Persons', 'Lug_Boot', 'Safety']]
Y = dataframe['Classification']

# Set the test size
test_size = 0.20 # Hyperparameter: Fraction of the dataset to use for
testing
seed = 7

# Split the dataset into test and train
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=test_size, random_state=seed)

# Create a K-Nearest Neighbors (K-NN) classifier
model = KNeighborsClassifier(n_neighbors=5, weights='uniform',
algorithm='auto')
# Hyperparameters:
# - n_neighbors: The number of nearest neighbors to consider when
making predictions. You can adjust this to control the model's
sensitivity to local patterns.
# - weights: Determines how the neighbors' contributions are weighted
(e.g., 'uniform' or 'distance'). You can choose the appropriate
weighting strategy.
# - algorithm: The algorithm used to compute the nearest neighbors
('auto', 'ball_tree', 'kd_tree', or 'brute'). You can choose the most
suitable algorithm based on your data size and structure.

# Train the model on the training data
model.fit(X_train, Y_train)

# Evaluate the accuracy
result = model.score(X_test, Y_test)
print("Accuracy: %.3f%%" % (result * 100.0))

Accuracy: 98.844%

```

5. Logistic Regression

- Sampling Technique - Train/Test Split (80:20)
- Classification Metrics - Accuracy

```
from pandas import read_csv
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/car_evaluation_classification.csv'
dataframe = read_csv(filename)
# Custom mapping for each feature
buying_price_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
maintenance_cost_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
lug_boot_mapping = {'big': 2, 'med': 1, 'small': 0}
safety_mapping = {'high': 2, 'med': 1, 'low': 0}

# Apply the custom mapping to each column
dataframe['Buying Price'] = dataframe['Buying
Price'].map(buying_price_mapping)
dataframe['Maintenance Cost'] = dataframe['Maintenance
Cost'].map(maintenance_cost_mapping)
dataframe['Lug_Boot'] = dataframe['Lug_Boot'].map(lug_boot_mapping)
dataframe['Safety'] = dataframe['Safety'].map(safety_mapping)

# features (X) and target (Y)
X = dataframe[['Buying Price', 'Maintenance Cost', 'Number of Doors',
'Number of Persons', 'Lug_Boot', 'Safety']]
Y = dataframe['Classification']

# Set the test size
test_size = 0.20 # Hyperparameter: Fraction of the dataset to use for
testing
seed = 7

# Split the dataset into test and train
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=test_size, random_state=seed)

# Create a Logistic Regression model
model = LogisticRegression(max_iter=200, solver='lbfgs', C=1.0)
# Hyperparameters:
# - max_iter: The maximum number of iterations for the solver to
converge. You can adjust this if the model does not converge.
# - solver: The algorithm to use for optimization ('lbfgs',
'liblinear', etc.). Choose an appropriate solver for your data and
problem.
# - C: Inverse of regularization strength. Smaller values increase
regularization. You can adjust this to control the trade-off between
```

fitting the data and preventing overfitting.

Train the model on the training data

```
model.fit(X_train, Y_train)
```

Evaluate the accuracy

```
result = model.score(X_test, Y_test)
```

```
print("Accuracy: %.3f%%" % (result * 100.0))
```

Accuracy: 89.017%

6. Multi-Layer Perceptron (MLP)

- Sampling Technique - Train/Test Split (80:20)
- Classification Metrics - Accuracy

```
from pandas import read_csv
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.neural_network import MLPClassifier
```

Load the dataset

```
filename = 'C:/Users/user/Desktop/ITD105
```

```
Files/CaseStudy1/Datasets/car_evaluation_classification.csv'
```

```
dataframe = read_csv(filename)
```

Custom mapping for each feature

```
buying_price_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
```

```
maintenance_cost_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
```

```
lug_boot_mapping = {'big': 2, 'med': 1, 'small': 0}
```

```
safety_mapping = {'high': 2, 'med': 1, 'low': 0}
```

Apply the custom mapping to each column

```
dataframe['Buying Price'] = dataframe['Buying  
Price'].map(buying_price_mapping)
```

```
dataframe['Maintenance Cost'] = dataframe['Maintenance  
Cost'].map(maintenance_cost_mapping)
```

```
dataframe['Lug_Boot'] = dataframe['Lug_Boot'].map(lug_boot_mapping)
```

```
dataframe['Safety'] = dataframe['Safety'].map(safety_mapping)
```

features (X) and target (Y)

```
X = dataframe[['Buying Price', 'Maintenance Cost', 'Number of Doors',  
'Number of Persons', 'Lug_Boot', 'Safety']]
```

```
Y = dataframe['Classification']
```

Set the test size

```
test_size = 0.20 # Hyperparameter: Fraction of the dataset to use for  
testing
```

```
seed = 7
```

Split the dataset into test and train

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,  
test_size=test_size, random_state=seed)
```

```

# Create an MLP-based model
model = MLPClassifier(hidden_layer_sizes=(65, 32), activation='relu',
solver='adam', max_iter=1000, random_state=seed)
# Hyperparameters:
# - hidden_layer_sizes: The number of neurons in each hidden layer.
You can customize the architecture by adjusting this parameter.
# - activation: The activation function used in the hidden layers
('relu', 'tanh', etc.). Choose the appropriate one for your problem.
# - solver: The algorithm for weight optimization ('adam', 'lbfgs',
etc.). Select the one that works best for your data.
# - max_iter: The maximum number of iterations for the solver to
converge. You can adjust this if the model does not converge.
# - random_state: The random seed for reproducibility. Set this to a
specific value for consistent results.

# Train the model
model.fit(X_train, Y_train)

# Evaluate the accuracy
result = model.score(X_test, Y_test)
print("Accuracy: %.3f%%" % (result * 100.0))

Accuracy: 99.422%

```

7. Perceptron

- Sampling Technique - Train/Test Split (80:20)
- Classification Metrics - Accuracy

```

from pandas import read_csv
from sklearn.model_selection import train_test_split
from sklearn.linear_model import Perceptron

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/car_evaluation_classification.csv'
dataframe = read_csv(filename)
# Custom mapping for each feature
buying_price_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
maintenance_cost_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
lug_boot_mapping = {'big': 2, 'med': 1, 'small': 0}
safety_mapping = {'high': 2, 'med': 1, 'low': 0}

# Apply the custom mapping to each column
dataframe['Buying Price'] = dataframe['Buying
Price'].map(buying_price_mapping)
dataframe['Maintenance Cost'] = dataframe['Maintenance
Cost'].map(maintenance_cost_mapping)
dataframe['Lug_Boot'] = dataframe['Lug_Boot'].map(lug_boot_mapping)

```



```

dataframe['Safety'] = dataframe['Safety'].map(safety_mapping)

# features (X) and target (Y)
X = dataframe[['Buying Price', 'Maintenance Cost', 'Number of Doors',
'Number of Persons', 'Lug_Boot', 'Safety']]
Y = dataframe['Classification']

# Set the test size
test_size = 0.20 # Hyperparameter: Fraction of the dataset to use for
testing
seed = 7

# Split the dataset into test and train
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=test_size, random_state=seed)

# Create a Perceptron classifier
model = Perceptron(max_iter=200, random_state=seed, eta0=1.0, tol=1e-
3)
# Hyperparameters:
# - max_iter: The maximum number of iterations for the solver to
converge. You can adjust this if the model does not converge.
# - random_state: The random seed for reproducibility. Set this to a
specific value for consistent results.
# - eta0: The initial learning rate. You can control the step size for
weight updates by adjusting this.
# - tol: The tolerance for stopping criterion. The model will stop
training when the change in the average loss is smaller than this
value.

# Train the model
model.fit(X_train, Y_train)

# Evaluate the accuracy
result = model.score(X_test, Y_test)
print("Accuracy: %.3f%%" % (result * 100.0))

Accuracy: 83.526%

```

8. Random Forest

- Sampling Technique - Train/Test Split (80:20)
- Classification Metrics - Accuracy

```

from pandas import read_csv
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/car_evaluation_classification.csv'

```

```

dataframe = read_csv(filename)
# Custom mapping for each feature
buying_price_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
maintenance_cost_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
lug_boot_mapping = {'big': 2, 'med': 1, 'small': 0}
safety_mapping = {'high': 2, 'med': 1, 'low': 0}

# Apply the custom mapping to each column
dataframe['Buying Price'] = dataframe['Buying
Price'].map(buying_price_mapping)
dataframe['Maintenance Cost'] = dataframe['Maintenance
Cost'].map(maintenance_cost_mapping)
dataframe['Lug_Boot'] = dataframe['Lug_Boot'].map(lug_boot_mapping)
dataframe['Safety'] = dataframe['Safety'].map(safety_mapping)

# features (X) and target (Y)
X = dataframe[['Buying Price', 'Maintenance Cost', 'Number of Doors',
'Number of Persons', 'Lug_Boot', 'Safety']]
Y = dataframe['Classification']

# Set the test size
test_size = 0.20 # Hyperparameter: Fraction of the dataset to use for
testing
seed = 7

# Split the dataset into test and train
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=test_size, random_state=seed)

# Create a Random Forest classifier
rfmodel = RandomForestClassifier(n_estimators=100, random_state=seed,
max_depth=None, min_samples_split=2, min_samples_leaf=1)
# Hyperparameters:
# - n_estimators: The number of decision trees in the random forest.
Adjust this to control the ensemble size.
# - random_state: The random seed for reproducibility. Set this to a
specific value for consistent results.
# - max_depth: The maximum depth of the decision trees. You can limit
tree depth to prevent overfitting.
# - min_samples_split: The minimum number of samples required to split
a node. Adjust this to control tree node splitting.
# - min_samples_leaf: The minimum number of samples required in a leaf
node. You can adjust this to control tree leaf size.

# Train the model
rfmodel.fit(X_train, Y_train)

# Evaluate the accuracy
result = rfmodel.score(X_test, Y_test)
print("Accuracy: %.3f%%" % (result * 100.0))

```

Accuracy: 99.422%

9. Support Vector Machines (SVM)

- Sampling Technique - Train/Test Split (80:20)
- Classification Metrics - Accuracy

```
from pandas import read_csv
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/car_evaluation_classification.csv'
dataframe = read_csv(filename)
# Custom mapping for each feature
buying_price_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
maintenance_cost_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
lug_boot_mapping = {'big': 2, 'med': 1, 'small': 0}
safety_mapping = {'high': 2, 'med': 1, 'low': 0}

# Apply the custom mapping to each column
dataframe['Buying Price'] = dataframe['Buying
Price'].map(buying_price_mapping)
dataframe['Maintenance Cost'] = dataframe['Maintenance
Cost'].map(maintenance_cost_mapping)
dataframe['Lug_Boot'] = dataframe['Lug_Boot'].map(lug_boot_mapping)
dataframe['Safety'] = dataframe['Safety'].map(safety_mapping)

# features (X) and target (Y)
X = dataframe[['Buying Price', 'Maintenance Cost', 'Number of Doors',
'Number of Persons', 'Lug_Boot', 'Safety']]
Y = dataframe['Classification']

# Set the test size
test_size = 0.20 # Hyperparameter: Fraction of the dataset to use for
testing
seed = 7

# Split the dataset into test and train
X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=test_size, random_state=seed)

# Create an SVM classifier
model = SVC(kernel='linear', C=1.0, random_state=seed)
# Hyperparameters:
# - kernel: The type of kernel to use ('linear', 'poly', 'rbf', etc.).
Choose the appropriate kernel for your problem.
# - C: The regularization parameter. Smaller values increase
regularization. You can adjust this to control the trade-off between
```

fitting the data and preventing overfitting.
- random_state: The random seed for reproducibility. Set this to a specific value for consistent results.

Train the model
model.fit(X_train, Y_train)

Evaluate the accuracy
result = model.score(X_test, Y_test)
print("Accuracy: %.3f%%" % (result * 100.0))

Accuracy: 89.595%

Algorithms Used for Regression

1. CART (Classification and Regression Trees)
2. Elastic Net
3. Gradient Boosting Machines (AdaBoost)
4. K-Nearest Neighbors (K-NN)
5. Lasso and Ridge Regression
6. Linear Regression
7. Multi-Layer Perceptron (MLP)
8. Random Forest
9. Support Vector Machines (SVM)

**When comparing models, a lower MAE is generally better.

1. CART (Classification and Regression Trees) - DecisionTree Regressor

- Sampling Technique = K-fold Cross Validation (k=10)
- Classification Metrics = MAE

```
import pandas as pd
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.tree import DecisionTreeRegressor

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/used_car_price.csv'
df = read_csv(filename)

# Features and target variable
X = df[['brand', 'year', 'mileage', 'fuel_type', 'transmission']]
Y = df['price']

# One-hot encoding for categorical variables (brand and fuel_type)
X = pd.get_dummies(X, columns=['brand', 'fuel_type', 'transmission'],
drop_first=True)
```

```

# Split the dataset into a 10-fold cross-validation
kfold = KFold(n_splits=10, random_state=None)

# Train the data on a Decision Tree Regressor
model = DecisionTreeRegressor(max_depth=None, min_samples_split=2,
min_samples_leaf=1, random_state=None)
# Hyperparameters:
# - max_depth: The maximum depth of the decision tree. You can limit
tree depth to prevent overfitting.
# - min_samples_split: The minimum number of samples required to split
an internal node. Adjust this to control node splitting.
# - min_samples_leaf: The minimum number of samples required in a leaf
node. You can adjust this to control leaf size.
# - random_state: The random seed for reproducibility. Set this to a
specific value for consistent results.

# Calculate the mean absolute error
scoring = 'neg_mean_absolute_error'
results = cross_val_score(model, X, Y, cv=kfold, scoring=scoring)
print("MAE: %.3f (%.3f)" % (-results.mean(), results.std()))

MAE: 191257.973 (37669.408)

```

2. Elastic Net

- Sampling Technique = K-fold Cross Validation (k=10)
- Classification Metrics = MAE

```

from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import ElasticNet

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/used_car_price.csv'
df = read_csv(filename)

# Features and target variable
X = df[['brand', 'year', 'mileage', 'fuel_type', 'transmission']]
Y = df['price']

# One-hot encoding for categorical variables (brand and fuel_type)
X = pd.get_dummies(X, columns=['brand', 'fuel_type', 'transmission'],
drop_first=True)

# Split the dataset into a 10-fold cross-validation
kfold = KFold(n_splits=10, random_state=None)

# Train the data on an Elastic Net model
model = ElasticNet(alpha=1.0, l1_ratio=0.5, max_iter=1000,

```

```

random_state=None)
# Hyperparameters:
# - alpha: The regularization parameter that controls the balance
between L1 (Lasso) and L2 (Ridge) penalties. Adjust this to control
the regularization strength.
# - l1_ratio: The mixing parameter for L1 and L2 penalties. A value of
0 corresponds to L2, 1 to L1, and values in between to combinations.
# - max_iter: The maximum number of iterations for the solver to
converge. You can adjust this if the model does not converge.
# - random_state: The random seed for reproducibility. Set this to a
specific value for consistent results.

# Calculate the mean absolute error
scoring = 'neg_mean_absolute_error'
results = cross_val_score(model, X, Y, cv=kfold, scoring=scoring)
print("MAE: %.3f (%.3f)" % (-results.mean(), results.std()))

MAE: 220326.393 (35703.212)

```

3. Gradient Boosting Machines (AdaBoost)

- Sampling Technique = K-fold Cross Validation (k=10)
- Classification Metrics = MAE

```

from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import AdaBoostRegressor

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/used_car_price.csv'
df = read_csv(filename)

# Features and target variable
X = df[['brand', 'year', 'mileage', 'fuel_type', 'transmission']]
Y = df['price']

# One-hot encoding for categorical variables (brand and fuel_type)
X = pd.get_dummies(X, columns=['brand', 'fuel_type', 'transmission'],
drop_first=True)

# Split the dataset into a 10-fold cross-validation
kfold = KFold(n_splits=10, random_state=None)

# Train the data on an AdaBoost Regressor
ada_model = AdaBoostRegressor(n_estimators=50, learning_rate=1.0,
random_state=None)
# Hyperparameters:
# - n_estimators: The number of weak regressors to combine in the
ensemble. You can adjust this to control the complexity of the

```

```

ensemble.
# - learning_rate: The contribution of each weak regressor to the
final prediction. You can adjust this to control the impact of
individual estimators.
# - random_state: The random seed for reproducibility. Set this to a
specific value for consistent results.

# Calculate the mean absolute error with AdaBoost
scoring = 'neg_mean_absolute_error'
ada_results = cross_val_score(ada_model, X, Y, cv=kfold,
scoring=scoring)
print("AdaBoost MAE: %.3f (%.3f)" % (-ada_results.mean(),
ada_results.std()))

AdaBoost MAE: 348184.901 (29348.249)

```

4. K-Nearest Neighbors (K-NN)

- Sampling Technique = K-fold Cross Validation (k=10)
- Classification Metrics = MAE

```

from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.neighbors import KNeighborsRegressor

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/used_car_price.csv'
df = read_csv(filename)

# Features and target variable
X = df[['brand', 'year', 'mileage', 'fuel_type', 'transmission']]
Y = df['price']

# One-hot encoding for categorical variables (brand and fuel_type)
X = pd.get_dummies(X, columns=['brand', 'fuel_type', 'transmission'],
drop_first=True)

# Split the dataset into a 10-fold cross-validation
kfold = KFold(n_splits=10, random_state=None)

# Train the data on a K-Nearest Neighbors Regressor
knn_model = KNeighborsRegressor(n_neighbors=5, weights='uniform',
algorithm='auto')
# Hyperparameters:
# - n_neighbors: The number of nearest neighbors to consider when
making predictions. You can adjust this to control the model's
sensitivity to local patterns.
# - weights: Determines how the neighbors' contributions are weighted
(e.g., 'uniform' or 'distance'). You can choose the appropriate

```

```
weighting strategy.  
# - algorithm: The algorithm used to compute the nearest neighbors  
('auto', 'ball_tree', 'kd_tree', or 'brute'). You can choose the most  
suitable algorithm based on your data size and structure.
```

```
# Calculate the mean absolute error with K-NN  
scoring = 'neg_mean_absolute_error'  
knn_results = cross_val_score(knn_model, X, Y, cv=kfold,  
scoring=scoring)  
print("K-NN MAE: %.3f (%.3f)" % (-knn_results.mean(),  
knn_results.std()))
```

K-NN MAE: 248792.473 (44887.474)

5. Lasso and Ridge Regression

- Sampling Technique = K-fold Cross Validation (k=10)
- Classification Metrics = MAE

```
from pandas import read_csv  
from sklearn.model_selection import KFold  
from sklearn.model_selection import cross_val_score  
from sklearn.linear_model import Lasso  
from sklearn.linear_model import Ridge  
  
# Load the dataset  
filename = 'C:/Users/user/Desktop/ITD105  
Files/CaseStudy1/Datasets/used_car_price.csv'  
df = read_csv(filename)  
  
# Features and target variable  
X = df[['brand', 'year', 'mileage', 'fuel_type', 'transmission']]  
Y = df['price']  
  
# One-hot encoding for categorical variables (brand and fuel_type)  
X = pd.get_dummies(X, columns=['brand', 'fuel_type', 'transmission'],  
drop_first=True)  
  
# Split the dataset into a 10-fold cross-validation  
kfold = KFold(n_splits=10, random_state=None)  
  
# Train the data on a Lasso Regression model  
lasso_model = Lasso(alpha=1.0, max_iter=1000, random_state=None)  
# Hyperparameters for Lasso:  
# - alpha: The regularization parameter that controls the strength of  
L1 regularization. Adjust this to control the level of sparsity in the  
model.  
# - max_iter: The maximum number of iterations for the solver to  
converge. You can adjust this if the model does not converge.  
# - random_state: The random seed for reproducibility. Set this to a  
specific value for consistent results.
```



```

# Calculate the mean absolute error with Lasso
scoring = 'neg_mean_absolute_error'
lasso_results = cross_val_score(lasso_model, X, Y, cv=kfold,
scoring=scoring)
print("Lasso MAE: %.3f (%.3f)" % (-lasso_results.mean(),
lasso_results.std()))

# Train the data on a Ridge Regression model
ridge_model = Ridge(alpha=1.0, max_iter=1000, random_state=None)
# Hyperparameters for Ridge:
# - alpha: The regularization parameter that controls the strength of
L2 regularization. Adjust this to control the strength of
regularization.
# - max_iter: The maximum number of iterations for the solver to
converge. You can adjust this if the model does not converge.
# - random_state: The random seed for reproducibility. Set this to a
specific value for consistent results.

# Calculate the mean absolute error with Ridge
ridge_results = cross_val_score(ridge_model, X, Y, cv=kfold,
scoring=scoring)
print("Ridge MAE: %.3f (%.3f)" % (-ridge_results.mean(),
ridge_results.std()))

Lasso MAE: 156500.518 (34501.360)
Ridge MAE: 163854.322 (31709.921)

```

6. Linear Regression

- Sampling Technique = K-fold Cross Validation (k=10)
- Classification Metrics = MAE

```

from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.linear_model import LinearRegression

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/used_car_price.csv'
df = read_csv(filename)

# Features and target variable
X = df[['brand', 'year', 'mileage', 'fuel_type', 'transmission']]
Y = df['price']

# One-hot encoding for categorical variables (brand and fuel_type)
X = pd.get_dummies(X, columns=['brand', 'fuel_type', 'transmission'],
drop_first=True)
# Split the dataset into a 10-fold cross-validation

```

```

kfold = KFold(n_splits=10, random_state=None)

# Train the data on a Linear Regression model
model = LinearRegression()

# Calculate the mean absolute error
scoring = 'neg_mean_absolute_error'
results = cross_val_score(model, X, Y, cv=kfold, scoring=scoring)
print("MAE: %.3f (%.3f)" % (-results.mean(), results.std()))

MAE: 156496.403 (34504.130)

```

7. Multi-Layer Perceptron (MLP)

- Sampling Technique = K-fold Cross Validation (k=10)
- Classification Metrics = MAE

```

from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.neural_network import MLPRegressor

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/used_car_price.csv'
df = read_csv(filename)

# Features and target variable
X = df[['brand', 'year', 'mileage', 'fuel_type', 'transmission']]
Y = df['price']

# One-hot encoding for categorical variables (brand and fuel_type)
X = pd.get_dummies(X, columns=['brand', 'fuel_type', 'transmission'],
drop_first=True)

# Split the dataset into a 10-fold cross-validation
kfold = KFold(n_splits=10, random_state=None)

# Train the data on an MLP Regressor with specified hyperparameters
mlp_model = MLPRegressor(
    hidden_layer_sizes=(100, 50), # Hyperparameter: Adjust the
    architecture as needed, specifying the number and size of hidden
    layers.
    activation='relu', # Hyperparameter: Choose an
    appropriate activation function ('identity', 'logistic', 'tanh',
    'relu', etc.).
    solver='adam', # Hyperparameter: Choose an
    optimization algorithm ('adam', 'lbfgs', 'sgd', etc.).
    learning_rate='constant', # Hyperparameter: Choose a learning
    rate schedule ('constant', 'invscaling', 'adaptive').
    max_iter=1000, # Hyperparameter: Adjust the maximum

```

```

number of iterations for training.
    random_state=50                # Hyperparameter: Set a random seed
for reproducibility.
)

```

```

# Calculate the mean absolute error with MLP
scoring = 'neg_mean_absolute_error'
mlp_results = cross_val_score(mlp_model, X, Y, cv=kfold,
scoring=scoring)
print("MLP MAE: %.3f (%.3f)" % (-mlp_results.mean(),
mlp_results.std()))

```

MLP MAE: 254932.805 (39260.489)

8. Random Forest

- Sampling Technique = K-fold Cross Validation (k=10)
- Classification Metrics = MAE

```

from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import RandomForestRegressor

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/used_car_price.csv'
df = read_csv(filename)

# Features and target variable
X = df[['brand', 'year', 'mileage', 'fuel_type', 'transmission']]
Y = df['price']

# One-hot encoding for categorical variables (brand and fuel_type)
X = pd.get_dummies(X, columns=['brand', 'fuel_type', 'transmission'],
drop_first=True)

# Split the dataset into a 10-fold cross-validation
kfold = KFold(n_splits=10, random_state=None)

# Train the data on a Random Forest Regressor with specified
hyperparameters
rf_model = RandomForestRegressor(
    n_estimators=100,          # Hyperparameter: The number of trees in
the forest. You can adjust this for ensemble size.
    max_depth=None,           # Hyperparameter: The maximum depth of
each tree. Adjust to control tree depth.
    min_samples_split=2,      # Hyperparameter: The minimum number of
samples required to split an internal node. Adjust to control node
splitting.
    min_samples_leaf=1,       # Hyperparameter: The minimum number of

```

```

samples required in a leaf node. Adjust to control leaf size.
    random_state=42          # Hyperparameter: Set a random seed for reproducibility.
)

```

```

# Calculate the mean absolute error with Random Forest
scoring = 'neg_mean_absolute_error'
rf_results = cross_val_score(rf_model, X, Y, cv=kfold,
scoring=scoring)
print("Random Forest MAE: %.3f (%.3f)" % (-rf_results.mean(),
rf_results.std()))

```

Random Forest MAE: 167836.651 (30518.495)

9. Support Vector Machines (SVM)

- Sampling Technique = K-fold Cross Validation (k=10)
- Classification Metrics = MAE

```

from pandas import read_csv
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
from sklearn.svm import SVR

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/used_car_price.csv'
df = read_csv(filename)

# Features and target variable
X = df[['brand', 'year', 'mileage', 'fuel_type', 'transmission']]
Y = df['price']

# One-hot encoding for categorical variables (brand and fuel_type)
X = pd.get_dummies(X, columns=['brand', 'fuel_type', 'transmission'],
drop_first=True)

# Split the dataset into a 10-fold cross-validation
kfold = KFold(n_splits=10, random_state=None)

# Train the data on a Support Vector Regressor (SVM) with specified hyperparameters
svm_model = SVR(
    kernel='rbf',          # Hyperparameter: The kernel function to use ('linear', 'poly', 'rbf', etc.).
    C=1.0,                # Hyperparameter: The regularization parameter. Adjust this to control the trade-off between margin width and error.
    epsilon=0.1,          # Hyperparameter: The epsilon-tube within which no penalty is associated with errors.
)

```

```

# Calculate the mean absolute error with SVM
scoring = 'neg_mean_absolute_error'
svm_results = cross_val_score(svm_model, X, Y, cv=kfold,
scoring=scoring)
print("SVM MAE: %.3f (%.3f)" % (-svm_results.mean(),
svm_results.std()))

SVM MAE: 241017.200 (48803.878)

import joblib # Import the joblib library
#Save the model to a file
#model_filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Models/classification_random_forest_model.pkl'
#joblib.dump(rfmodel, model_filename)

import joblib

# Load the saved Random Forest model
model_filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Models/classification_random_forest_model.pkl'
loaded_model = joblib.load(model_filename)

# Define sample input data
#Buying Price    Maintenance Cost Number of Doors Number of Persons
    Lug_Boot    Safety
sample_input = [['1', '1', '2', '1', '2', '2',]]

# Make predictions using the loaded model
predictions = loaded_model.predict(sample_input)

# Define messages based on the predicted class
if predictions[0] == 1:
    print("Accepted")
else:
    print("Unaccepted")

Unaccepted

C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-
packages\sklearn\base.py:420: UserWarning: X does not have valid
feature names, but RandomForestClassifier was fitted with feature
names
    warnings.warn(

```

Comparing ML Algorithms

Classification

```
#Split and Train Test
#Confusion Matrix and Classification Report

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neural_network import MLPClassifier
from sklearn.linear_model import Perceptron
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/car_evaluation_classification.csv'
dataframe = read_csv(filename)
# Custom mapping for each feature
buying_price_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
maintenance_cost_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
lug_boot_mapping = {'big': 2, 'med': 1, 'small': 0}
safety_mapping = {'high': 2, 'med': 1, 'low': 0}

# Apply the custom mapping to each column
dataframe['Buying Price'] = dataframe['Buying
Price'].map(buying_price_mapping)
dataframe['Maintenance Cost'] = dataframe['Maintenance
Cost'].map(maintenance_cost_mapping)
dataframe['Lug_Boot'] = dataframe['Lug_Boot'].map(lug_boot_mapping)
dataframe['Safety'] = dataframe['Safety'].map(safety_mapping)

# features (X) and target (Y)
X = dataframe[['Buying Price', 'Maintenance Cost', 'Number of Doors',
'Number of Persons', 'Lug_Boot', 'Safety']]
y = dataframe['Classification']

# Split the data into a training and testing set
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Initialize classifiers
```

```

classifiers = {
    "CART": DecisionTreeClassifier(),
    "Naive Bayes": GaussianNB(),
    "AdaBoost": AdaBoostClassifier(),
    "K-NN": KNeighborsClassifier(),
    "Logistic Regression": LogisticRegression(),
    "MLP": MLPClassifier(),
    "Perceptron": Perceptron(),
    "Random Forest": RandomForestClassifier(),
    "SVM": SVC(),
}

# Iterate through classifiers and evaluate
for name, clf in classifiers.items():
    clf.fit(X_train, y_train)
    y_pred = clf.predict(X_test)

    # Evaluate and print confusion matrix and classification report
    print(f"Classifier: {name}")
    print("Confusion Matrix:")
    print(confusion_matrix(y_test, y_pred))
    print("Classification Report:")
    print(classification_report(y_test, y_pred))
    print("\n")

```

```

Classifier: CART
Confusion Matrix:
[[234  1]
 [ 2 109]]
Classification Report:

```

	precision	recall	f1-score	support
0	0.99	1.00	0.99	235
1	0.99	0.98	0.99	111
accuracy			0.99	346
macro avg	0.99	0.99	0.99	346
weighted avg	0.99	0.99	0.99	346

```

Classifier: Naive Bayes
Confusion Matrix:
[[220 15]
 [ 19 92]]
Classification Report:

```

	precision	recall	f1-score	support
0	0.92	0.94	0.93	235
1	0.86	0.83	0.84	111

accuracy			0.90	346
macro avg	0.89	0.88	0.89	346
weighted avg	0.90	0.90	0.90	346

Classifier: AdaBoost

Confusion Matrix:

[[228 7]

[11 100]]

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.97	0.96	235
1	0.93	0.90	0.92	111

accuracy			0.95	346
macro avg	0.94	0.94	0.94	346
weighted avg	0.95	0.95	0.95	346

Classifier: K-NN

Confusion Matrix:

[[235 0]

[6 105]]

Classification Report:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	235
1	1.00	0.95	0.97	111

accuracy			0.98	346
macro avg	0.99	0.97	0.98	346
weighted avg	0.98	0.98	0.98	346

Classifier: Logistic Regression

Confusion Matrix:

[[220 15]

[26 85]]

Classification Report:

	precision	recall	f1-score	support
0	0.89	0.94	0.91	235
1	0.85	0.77	0.81	111

accuracy			0.88	346
macro avg	0.87	0.85	0.86	346

weighted avg	0.88	0.88	0.88	346
--------------	------	------	------	-----

```
C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-  
packages\sklearn\normal_network\_multilayer_perceptron.py:684:  
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200)  
reached and the optimization hasn't converged yet.  
warnings.warn(  

```

Classifier: MLP

Confusion Matrix:

```
[[231  4]
```

```
 [ 11 100]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.98	0.97	235
1	0.96	0.90	0.93	111
accuracy			0.96	346
macro avg	0.96	0.94	0.95	346
weighted avg	0.96	0.96	0.96	346

Classifier: Perceptron

Confusion Matrix:

```
[[190  45]
```

```
 [  5 106]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.97	0.81	0.88	235
1	0.70	0.95	0.81	111
accuracy			0.86	346
macro avg	0.84	0.88	0.85	346
weighted avg	0.89	0.86	0.86	346

Classifier: Random Forest

Confusion Matrix:

```
[[235  0]
```

```
 [  2 109]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.99	1.00	1.00	235

1	1.00	0.98	0.99	111
accuracy			0.99	346
macro avg	1.00	0.99	0.99	346
weighted avg	0.99	0.99	0.99	346

Classifier: SVM

Confusion Matrix:

```
[[233  2]
 [ 10 101]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.96	0.99	0.97	235
1	0.98	0.91	0.94	111
accuracy			0.97	346
macro avg	0.97	0.95	0.96	346
weighted avg	0.97	0.97	0.96	346

#K-Fold Cross Variation
#Classification Accuracy

Import necessary libraries

```
from sklearn.model_selection import cross_val_score
from sklearn.metrics import accuracy_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.ensemble import AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neural_network import MLPClassifier
from sklearn.linear_model import Perceptron
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
```

Load the dataset

```
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/car_evaluation_classification.csv'
dataframe = read_csv(filename)
```

Custom mapping for each feature

```
buying_price_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
maintenance_cost_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
lug_boot_mapping = {'big': 2, 'med': 1, 'small': 0}
safety_mapping = {'high': 2, 'med': 1, 'low': 0}
```

```

# Apply the custom mapping to each column
dataframe['Buying Price'] = dataframe['Buying
Price'].map(buying_price_mapping)
dataframe['Maintenance Cost'] = dataframe['Maintenance
Cost'].map(maintenance_cost_mapping)
dataframe['Lug_Boot'] = dataframe['Lug_Boot'].map(lug_boot_mapping)
dataframe['Safety'] = dataframe['Safety'].map(safety_mapping)

# Define features (X) and target (y)
X = dataframe[['Buying Price', 'Maintenance Cost', 'Number of Doors',
'Number of Persons', 'Lug_Boot', 'Safety']]
y = dataframe['Classification']

# Define a list of machine learning algorithms
algorithms = [
    ("CART", DecisionTreeClassifier()),
    ("Naive Bayes", GaussianNB()),
    ("AdaBoost", AdaBoostClassifier()),
    ("K-NN", KNeighborsClassifier()),
    ("Logistic Regression", LogisticRegression()),
    ("MLP", MLPClassifier()),
    ("Perceptron", Perceptron()),
    ("Random Forest", RandomForestClassifier()),
    ("SVM", SVC())
]

# Iterate through the algorithms and perform K-fold Cross Validation
for name, model in algorithms:
    scores = cross_val_score(model, X, y, cv=10, scoring='accuracy')
# K-fold Cross Validation with k=10
    print(f"{name}:")
    print(f"Mean Accuracy: {scores.mean()}")
    print(f"Standard Deviation: {scores.std()}\n")

```

CART:
Mean Accuracy: 0.8473753192633419
Standard Deviation: 0.11940050981103918

Naive Bayes:
Mean Accuracy: 0.8993816373168437
Standard Deviation: 0.06628282497867485

AdaBoost:
Mean Accuracy: 0.917898911143971
Standard Deviation: 0.06281304813811818

K-NN:
Mean Accuracy: 0.8941457185105526
Standard Deviation: 0.0398902191529915

Logistic Regression:
Mean Accuracy: 0.8709067078908456
Standard Deviation: 0.06944590627065844

C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:684:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

warnings.warn(

C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:684:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

warnings.warn(

C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:684:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

warnings.warn(

C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:684:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

warnings.warn(

C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:684:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

warnings.warn(

C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:684:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

warnings.warn(

C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:684:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

warnings.warn(

C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:684:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

warnings.warn(

C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\neural_network_multilayer_perceptron.py:684:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200) reached and the optimization hasn't converged yet.

warnings.warn(

```
C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-  
packages\sklearn\normal_network\_multilayer_perceptron.py:684:  
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200)  
reached and the optimization hasn't converged yet.  
warnings.warn(  

```

MLP:

Mean Accuracy: 0.9375352870009408

Standard Deviation: 0.044445170561649996

Perceptron:

Mean Accuracy: 0.8270668100551151

Standard Deviation: 0.1008261198126333

Random Forest:

Mean Accuracy: 0.9063314961688398

Standard Deviation: 0.08069074088344796

SVM:

Mean Accuracy: 0.9288546847694583

Standard Deviation: 0.04578712545956918

```
from sklearn.model_selection import RepeatedKFold  
from sklearn.model_selection import cross_val_score  
from sklearn.metrics import log_loss  
from sklearn.tree import DecisionTreeClassifier  
from sklearn.naive_bayes import GaussianNB  
from sklearn.ensemble import AdaBoostClassifier  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.linear_model import LogisticRegression  
from sklearn.neural_network import MLPClassifier  
from sklearn.ensemble import RandomForestClassifier  
from sklearn.svm import SVC
```

```
# Define your dataset and features (X) and target (y) here
```

```
# Load and preprocess your data
```

```
# Define resampling technique: Repeated Random Train-Test splits
```

```
n_splits = 5 # Number of splits
```

```
n_repeats = 10 # Number of repetitions
```

```
rkf = RepeatedKFold(n_splits=n_splits, n_repeats=n_repeats,  
random_state=42)
```

```
# Define classification algorithms (excluding Perceptron)
```

```
classifiers = {  
    'CART': DecisionTreeClassifier(),  
    'Gaussian NB': GaussianNB(),  
    'AdaBoost': AdaBoostClassifier(),  
    'K-NN': KNeighborsClassifier(),  

```

```

'Logistic Regression': LogisticRegression(),
'MLP': MLPClassifier(), # Adjust max_iter if needed
'Random Forest': RandomForestClassifier(),
'SVM': SVC(probability=True) # Use probability=True for log loss
}

# Evaluate each classifier using log loss
for name, clf in classifiers.items():
    try:
        logloss_scores = -cross_val_score(clf, X, y, cv=rkf,
scoring='neg_log_loss')
        print(f'{name} Log Loss: {logloss_scores.mean():.4f} (±
{logloss_scores.std():.4f})')
    except Exception as e:
        print(f'{name} Log Loss: N/A (Exception: {str(e)})')

CART Log Loss: 0.3400 (±0.1841)
Gaussian NB Log Loss: 0.2055 (±0.0157)
AdaBoost Log Loss: 0.4945 (±0.0083)
K-NN Log Loss: 0.1165 (±0.0218)
Logistic Regression Log Loss: 0.2463 (±0.0234)

C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-
packages\sklearn\normalization\_multilayer_perceptron.py:684:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (1000)
reached and the optimization hasn't converged yet.
warnings.warn(
C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-
packages\sklearn\normalization\_multilayer_perceptron.py:684:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (1000)
reached and the optimization hasn't converged yet.
warnings.warn(

MLP Log Loss: 0.0633 (±0.0165)
Random Forest Log Loss: 0.0623 (±0.0070)
SVM Log Loss: 0.0904 (±0.0171)

#Random Forest Hypertuning

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/car_evaluation_classification.csv'
dataframe = read_csv(filename)
# Custom mapping for each feature
buying_price_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}
maintenance_cost_mapping = {'vhigh': 3, 'high': 2, 'med': 1, 'low': 0}

```

```

lug_boot_mapping = {'big': 2, 'med': 1, 'small': 0}
safety_mapping = {'high': 2, 'med': 1, 'low': 0}

# Apply the custom mapping to each column
dataframe['Buying Price'] = dataframe['Buying
Price'].map(buying_price_mapping)
dataframe['Maintenance Cost'] = dataframe['Maintenance
Cost'].map(maintenance_cost_mapping)
dataframe['Lug_Boot'] = dataframe['Lug_Boot'].map(lug_boot_mapping)
dataframe['Safety'] = dataframe['Safety'].map(safety_mapping)

# Define features (X) and target (y)
X = dataframe[['Buying Price', 'Maintenance Cost', 'Number of Doors',
'Number of Persons', 'Lug_Boot', 'Safety']]
y = dataframe['Classification']

# Split the dataset into training and testing sets (80% training, 20%
testing)
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Define a list of hyperparameters to tune
n_estimators = [50, 100, 200] # Number of trees in the forest
max_depth = [None, 10, 20] # Maximum depth of each tree
min_samples_split = [2, 5, 10] # Minimum number of samples required
to split an internal node

best_accuracy = 0.0
best_model = None

# Perform hyperparameter tuning
for n in n_estimators:
    for depth in max_depth:
        for min_samples in min_samples_split:
            # Create a Random Forest classifier with the current
hyperparameters
            rf = RandomForestClassifier(n_estimators=n,
max_depth=depth, min_samples_split=min_samples, random_state=42)

            # Train the model on the training data
            rf.fit(X_train, y_train)

            # Make predictions on the test data
            y_pred = rf.predict(X_test)

            # Calculate accuracy
            accuracy = accuracy_score(y_test, y_pred)

            # Check if this model achieved the best accuracy
            if accuracy > best_accuracy:

```

```

        best_accuracy = accuracy
        best_model = rf

        print(f"n_estimators: {n}, max_depth: {depth},
min_samples_split: {min_samples}, Accuracy: {accuracy:.3f}")

# Print the best hyperparameters and accuracy
print("\nBest Hyperparameters:")
print(f"n_estimators: {best_model.n_estimators}, max_depth:
{best_model.max_depth}, "
      f"min_samples_split: {best_model.min_samples_split}, Accuracy:
{best_accuracy:.3f}")

n_estimators: 50, max_depth: None, min_samples_split: 2, Accuracy:
0.997
n_estimators: 50, max_depth: None, min_samples_split: 5, Accuracy:
0.994
n_estimators: 50, max_depth: None, min_samples_split: 10, Accuracy:
0.991
n_estimators: 50, max_depth: 10, min_samples_split: 2, Accuracy: 0.997
n_estimators: 50, max_depth: 10, min_samples_split: 5, Accuracy: 0.997
n_estimators: 50, max_depth: 10, min_samples_split: 10, Accuracy:
0.991
n_estimators: 50, max_depth: 20, min_samples_split: 2, Accuracy: 0.997
n_estimators: 50, max_depth: 20, min_samples_split: 5, Accuracy: 0.994
n_estimators: 50, max_depth: 20, min_samples_split: 10, Accuracy:
0.991
n_estimators: 100, max_depth: None, min_samples_split: 2, Accuracy:
0.997
n_estimators: 100, max_depth: None, min_samples_split: 5, Accuracy:
0.997
n_estimators: 100, max_depth: None, min_samples_split: 10, Accuracy:
0.994
n_estimators: 100, max_depth: 10, min_samples_split: 2, Accuracy:
0.997
n_estimators: 100, max_depth: 10, min_samples_split: 5, Accuracy:
0.997
n_estimators: 100, max_depth: 10, min_samples_split: 10, Accuracy:
0.994
n_estimators: 100, max_depth: 20, min_samples_split: 2, Accuracy:
0.997
n_estimators: 100, max_depth: 20, min_samples_split: 5, Accuracy:
0.997
n_estimators: 100, max_depth: 20, min_samples_split: 10, Accuracy:
0.994
n_estimators: 200, max_depth: None, min_samples_split: 2, Accuracy:
0.997
n_estimators: 200, max_depth: None, min_samples_split: 5, Accuracy:
0.997
n_estimators: 200, max_depth: None, min_samples_split: 10, Accuracy:

```



```
0.994
n_estimators: 200, max_depth: 10, min_samples_split: 2, Accuracy:
0.997
n_estimators: 200, max_depth: 10, min_samples_split: 5, Accuracy:
0.994
n_estimators: 200, max_depth: 10, min_samples_split: 10, Accuracy:
0.994
n_estimators: 200, max_depth: 20, min_samples_split: 2, Accuracy:
0.997
n_estimators: 200, max_depth: 20, min_samples_split: 5, Accuracy:
0.997
n_estimators: 200, max_depth: 20, min_samples_split: 10, Accuracy:
0.994

Best Hyperparameters:
n_estimators: 50, max_depth: None, min_samples_split: 2, Accuracy:
0.997
```

Regression

```
#Split into Train and Test Sets
#Mean Absolute Error (MAE)

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import ElasticNet, Lasso, Ridge,
LinearRegression
from sklearn.ensemble import AdaBoostRegressor, RandomForestRegressor,
GradientBoostingRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.metrics import mean_absolute_error

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/used_car_price.csv'
df = read_csv(filename)

# Features and target variable
X = df[['brand', 'year', 'mileage', 'fuel_type', 'transmission']]
Y = df['price']

# One-hot encoding for categorical variables (brand and fuel_type)
X = pd.get_dummies(X, columns=['brand', 'fuel_type', 'transmission'],
drop_first=True)

# Split the data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, Y,
test_size=0.2, random_state=42)
```

```
# Initialize a dictionary to store the results
results = {}
```

```
# Define and evaluate each regression algorithm
regressors = {
```

```
    'CART': DecisionTreeRegressor(),
    'Elastic Net': ElasticNet(),
    'AdaBoost': AdaBoostRegressor(),
    'K-NN': KNeighborsRegressor(),
    'Lasso Regression': Lasso(),
    'Ridge Regression': Ridge(),
    'Linear Regression': LinearRegression(),
    'Random Forest': RandomForestRegressor(),
    'Gradient Boosting': GradientBoostingRegressor(),
    'MLP': MLPRegressor()
}
```

```
for name, regressor in regressors.items():
    regressor.fit(X_train, y_train)
    y_pred = regressor.predict(X_test)
    mae = mean_absolute_error(y_test, y_pred)
    results[name] = mae
```

```
# Print the MAE for each algorithm
```

```
for name, mae in results.items():
    print(f'{name}: Mean Absolute Error = {mae:.2f}')
```

```
CART: Mean Absolute Error = 196970.82
Elastic Net: Mean Absolute Error = 265543.15
AdaBoost: Mean Absolute Error = 306335.08
K-NN: Mean Absolute Error = 280099.82
Lasso Regression: Mean Absolute Error = 199079.56
Ridge Regression: Mean Absolute Error = 209465.26
Linear Regression: Mean Absolute Error = 199070.78
Random Forest: Mean Absolute Error = 180124.27
Gradient Boosting: Mean Absolute Error = 205418.92
MLP: Mean Absolute Error = 360855.99
```

```
C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-
packages\sklearn\neural_network\multilayer_perceptron.py:684:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200)
reached and the optimization hasn't converged yet.
    warnings.warn(
```

```
#K-Fold Cross Variation
```

```
#R^2
```

```
import pandas as pd
```

```

from sklearn.model_selection import cross_val_score, KFold
from sklearn.tree import DecisionTreeRegressor
from sklearn.linear_model import ElasticNet, Lasso, Ridge,
LinearRegression
from sklearn.ensemble import AdaBoostRegressor, RandomForestRegressor,
GradientBoostingRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.metrics import r2_score

# Load the dataset
filename = 'C:/Users/user/Desktop/ITD105
Files/CaseStudy1/Datasets/used_car_price.csv'
df = pd.read_csv(filename)

# Features and target variable
X = df[['brand', 'year', 'mileage', 'fuel_type', 'transmission']]
y = df['price']

# One-hot encoding for categorical variables (brand and fuel_type)
X = pd.get_dummies(X, columns=['brand', 'fuel_type', 'transmission'],
drop_first=True)

# Initialize a dictionary to store the results
results = {}

# Define and evaluate each regression algorithm using K-fold Cross
Validation
regressors = {
    'CART': DecisionTreeRegressor(),
    'Elastic Net': ElasticNet(),
    'AdaBoost': AdaBoostRegressor(),
    'K-NN': KNeighborsRegressor(),
    'Lasso Regression': Lasso(),
    'Ridge Regression': Ridge(),
    'Linear Regression': LinearRegression(),
    'Random Forest': RandomForestRegressor(),
    'Gradient Boosting': GradientBoostingRegressor(),
    'MLP': MLPRegressor()
}

k_fold = KFold(n_splits=5, shuffle=True, random_state=42)

for name, regressor in regressors.items():
    scores = cross_val_score(regressor, X, y, cv=k_fold, scoring='r2')
    mean_r2 = scores.mean()
    results[name] = mean_r2

# Print the R-squared (R^2) for each algorithm

```

```
for name, r2 in results.items():  
    print(f'{name}: Mean R-squared (R^2) = {r2:.2f}')
```

```
C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-  
packages\sklearn\neural_network\_multilayer_perceptron.py:684:  
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200)  
reached and the optimization hasn't converged yet.
```

```
warnings.warn(  

```

```
C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-  
packages\sklearn\neural_network\_multilayer_perceptron.py:684:  
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200)  
reached and the optimization hasn't converged yet.
```

```
warnings.warn(  

```

```
C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-  
packages\sklearn\neural_network\_multilayer_perceptron.py:684:  
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200)  
reached and the optimization hasn't converged yet.
```

```
warnings.warn(  

```

```
C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-  
packages\sklearn\neural_network\_multilayer_perceptron.py:684:  
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200)  
reached and the optimization hasn't converged yet.
```

```
warnings.warn(  

```

```
CART: Mean R-squared (R^2) = -0.53
```

```
Elastic Net: Mean R-squared (R^2) = 0.15
```

```
AdaBoost: Mean R-squared (R^2) = -0.60
```

```
K-NN: Mean R-squared (R^2) = -0.08
```

```
Lasso Regression: Mean R-squared (R^2) = 0.53
```

```
Ridge Regression: Mean R-squared (R^2) = 0.53
```

```
Linear Regression: Mean R-squared (R^2) = 0.53
```

```
Random Forest: Mean R-squared (R^2) = 0.32
```

```
Gradient Boosting: Mean R-squared (R^2) = 0.46
```

```
MLP: Mean R-squared (R^2) = -0.38
```

```
C:\Users\user\AppData\Local\Programs\Python\Python310\lib\site-  
packages\sklearn\neural_network\_multilayer_perceptron.py:684:  
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (200)  
reached and the optimization hasn't converged yet.
```

```
warnings.warn(  

```

```
#Random Forest MAE Hypertuning
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.ensemble import RandomForestRegressor
```

```
from sklearn.metrics import mean_absolute_error
```

```
# Load the dataset
```

```
filename = 'C:/Users/user/Desktop/ITD105
```

```
Files/CaseStudy1/Datasets/used_car_price.csv'
```

```

df = pd.read_csv(filename)

# Features and target variable
X = df[['brand', 'year', 'mileage', 'fuel_type', 'transmission']]
Y = df['price']

# One-hot encoding for categorical variables (brand and fuel_type)
X = pd.get_dummies(X, columns=['brand', 'fuel_type', 'transmission'],
drop_first=True)

# Define the hyperparameters to be tuned
n_estimators_list = [50, 100, 200]
max_depth_list = [None, 10, 20]
min_samples_split_list = [2, 5, 10]

# Initialize a dictionary to store the results
results = {}

# Loop through different hyperparameters
for n_estimators in n_estimators_list:
    for max_depth in max_depth_list:
        for min_samples_split in min_samples_split_list:
            # Split the data into training and testing sets
            X_train, X_test, y_train, y_test = train_test_split(X, Y,
test_size=0.2, random_state=42)

            # Initialize and train the Random Forest model with the
current hyperparameters
            model = RandomForestRegressor(n_estimators=n_estimators,
max_depth=max_depth, min_samples_split=min_samples_split)
            model.fit(X_train, y_train)

            # Make predictions
            y_pred = model.predict(X_test)

            # Calculate MAE
            mae = mean_absolute_error(y_test, y_pred)

            # Store the MAE in the results dictionary
            results[f'{n_estimators} {max_depth} {min_samples_split}']
= mae

# Print the MAE for each set of hyperparameters
for params, mae in results.items():
    print(f'Hyperparameters: {params}, MAE = {mae:.2f}')

Hyperparameters: 50 None 2, MAE = 182228.24
Hyperparameters: 50 None 5, MAE = 183944.92
Hyperparameters: 50 None 10, MAE = 191451.29
Hyperparameters: 50 10 2, MAE = 203781.52

```

Hyperparameters: 50 10 5, MAE = 205453.82
Hyperparameters: 50 10 10, MAE = 208326.99
Hyperparameters: 50 20 2, MAE = 184014.74
Hyperparameters: 50 20 5, MAE = 187392.50
Hyperparameters: 50 20 10, MAE = 190737.77
Hyperparameters: 100 None 2, MAE = 181267.83
Hyperparameters: 100 None 5, MAE = 185898.92
Hyperparameters: 100 None 10, MAE = 189268.17
Hyperparameters: 100 10 2, MAE = 202844.31
Hyperparameters: 100 10 5, MAE = 204037.01
Hyperparameters: 100 10 10, MAE = 207474.02
Hyperparameters: 100 20 2, MAE = 184056.43
Hyperparameters: 100 20 5, MAE = 185613.13
Hyperparameters: 100 20 10, MAE = 190086.56
Hyperparameters: 200 None 2, MAE = 182771.66
Hyperparameters: 200 None 5, MAE = 184627.20
Hyperparameters: 200 None 10, MAE = 189835.21
Hyperparameters: 200 10 2, MAE = 203040.85
Hyperparameters: 200 10 5, MAE = 204136.29
Hyperparameters: 200 10 10, MAE = 205841.15
Hyperparameters: 200 20 2, MAE = 180738.81
Hyperparameters: 200 20 5, MAE = 185088.30
Hyperparameters: 200 20 10, MAE = 189119.05